

# Forward Pass Project Results

0. Batch profiling results for basic forward convolution kernel:

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	3.90226 ms	5.1347 ms	1.276 s	.86
1000	11.0889 ms	48.1185 ms	11.516 s	.886
10000	23.2205 ms	146.022 ms	1 min 40.687 s	.8714

## **1. Weight matrix in constant memory**

a. Which optimization did you choose to implement? Explain why did you choose that optimization technique.

I chose to implement a weight matrix in constant memory because I thought it would make the kernel more memory efficient since all of the operations can be stored and reused when needed from one single location.

b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

This optimization works by utilizing a matrix placed in constant memory to act as a filter over the various input images in the fashion data set. It is very memory efficient as constant memory is faster to access thus reducing the bandwidth requirements. I thought that this optimization would increase performance of the forward convolution since the data we would be dealing with would be anywhere from 100 to 10000 in terms of batch size, so the convolution would be expedited if there were less memory transactions required due to the matrix allowing better data accessibility. This optimization synergizes well with tiling and other shared memory techniques.

c. Batch profiling results:

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	.226164 ms	.839785 ms	0m 1.220s	.86
1000	2.10433 ms	8.13764 ms	0m 10.177s	.886
10000	20.6525 ms	81.62 ms	1m 40.276s	.8714

- d. Was implementing this optimization successful in improving performance? Why or why not?

*nsys Profile data:*

```
Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)
```

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
65.0	1171453601	6	195242266.8	23077	639816163	cudaMemcpy
29.1	524963137	8	65620442.1	89776	519520595	cudaMalloc
5.7	102436821	8	12804602.6	4234	81315873	cudaDeviceSynchronize
0.2	2983817	8	372877.1	69333	998325	cudaFree
0.0	314382	6	52397.0	22328	167585	cudaLaunchKernel
0.0	165118	2	82559.0	79492	85626	cudaMemcpyToSymbol

```
Generating CUDA Kernel Statistics...
Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)
```

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	102262558	2	51131279.0	20955574	81306984	conv_forward_kernel
0.0	2912	2	1456.0	1440	1472	do_not_remove_this_kernel
0.0	2720	2	1360.0	1344	1376	prefn_marker_kernel

```
CUDA Memory Operation Statistics (nanoseconds)
```

Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
93.5	1098382962	2	545151481.0	451516419	638786543	[CUDA memcpy DtoH]
6.5	75247344	6	12541224.0	1536	40845624	[CUDA memcpy HtoD]

```
CUDA Memory Operation Statistics (KiB)
```


Total	Operations	Average	Minimum	Maximum	Name
1722508.0	2	861250.0	722508.000	1868968.0	[CUDA memcpy DtoH]
538919.0	6	89819.0	0.004	288966.0	[CUDA memcpy HtoD]

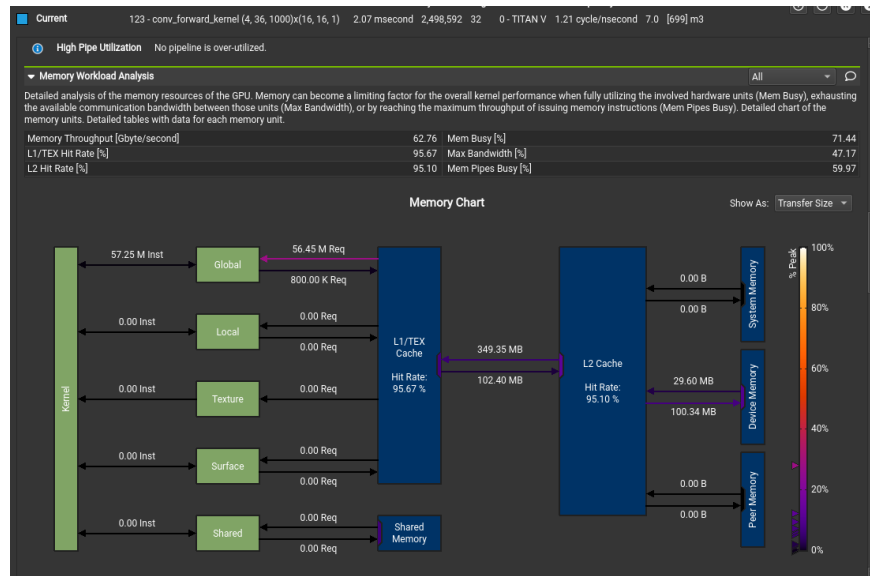
```
Generating Operating System Runtime API Statistics...
Operating System Runtime API Statistics (nanoseconds)
```

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
---------	------------	-------	---------	---------	---------	------

The time (%) displayed by the CUDA Kernel Statistics show that 65.0% of time was taken by the cudaMemcpy. Compared to the figure of 71.6% of time needed for cudaMemcpy from the baseline code measured for Milestone 2, this means that there was a reduction of over 6% in time. This can be attributed to successfully implementing the weighted matrix in constant memory which is why cudaMalloc increases from 15.7% to 29.1%. We can conclude that this optimization was successful due to lowered OP times (~102ms) and the reduction in time taken by cudaMemcpy.

*Nsight-Compute Profile data:*





Based on the Nsight-Compute data, we can observe that there were high amounts of SM and Memory utilization at figures of approximately 85% and 70% respectively which are relatively high figures compared to 75% and 60% respectively from Milestone 2, but it is important to look at the memory workload analysis. Judging by how the L1 and L2 Cache have hit rates of approximately 95%, this means that the optimization of implementing a weighted matrix in constant memory is memory efficient and we were successful in increasing performance.

## e. Modified code:

```

5 #define TILE_WIDTH 16
6
7 __constant__ float matrix[10000];
8 __global__ void conv_forward_kernel(float *output, const float *input, const float *mask, const int Batch, const int Map_out, const int Channel, con
9 {
10     /*
11     Modify this function to implement the forward pass described in Chapter 16.
12     We have added an additional dimension to the tensors to support an entire mini-batch
13     The goal here is to be correct AND fast.
14
15     Function parameter definitions:
16     output - output
17     input - input
18     mask - convolution kernel
19     Batch - batch_size (number of images in x)
20     Map_out - number of output feature maps
21     Channel - number of input feature maps
22     Height - input height dimension
23     Width - input width dimension
24     K - kernel height and width (K x K)
25     */
26
27     const int Height_out = Height - K + 1;
28     const int Width_out = Width - K + 1;
29     const int W_grd = (Width_out + TILE_WIDTH - 1) / TILE_WIDTH;
30
31     // We have some nice #defs for you below to simplify indexing. Feel free to use them, or create your own.
32     // An example use of these macros:
33     // float a = in_4d(0,0,0,0);
34     // out_4d(0,0,0,0) = a;
35
36     #define out_4d(l3, l2, l1, l0) output[(l3) * (Map_out * Height_out * Width_out) + (l2) * (Height_out * Width_out) + (l1) * (Width_out) + l0]
37     #define in_4d(l3, l2, l1, l0) input[(l3) * (Channel * Height * Width) + (l2) * (Height * Width) + (l1) * (Width) + l0]
38     #define mask_4d(l3, l2, l1, l0) matrix[(l3) * (Channel * K * K) + (l2) * (K * K) + (l1) * (K) + l0]
39
40     // Insert your GPU convolution kernel code here
41     int m = blockIdx.x;
42     int h = (blockIdx.y / W_grd) * TILE_WIDTH + threadIdx.y;
43     int w = (blockIdx.y % W_grd) * TILE_WIDTH + threadIdx.x;
44     int b = blockIdx.z;
45
46     float acc = 0.0f;
47     for(int c = 0; c < Channel; c++) { // sum over all input channels
48         for(int p = 0; p < K; p++) { // loop over KxK filter
49             for(int q = 0; q < K; q++)
50                 acc += in_4d(b, c, h + p, w + q) * mask_4d(m, c, p, q);
51         }
52     }
53 }

```

Code also modified in conv\_forward\_gpu\_prolog.

```

63 __host__ void GPUInterface::conv_forward_gpu_prolog(const float *host_output, const float *host_input, const float
**device_mask_ptr, const int Batch, const int Map_out, const int Channel, const int Height, const int Width, cons
64 {
65     // Allocate memory and copy over the relevant data structures to the GPU
66
67     // We pass double pointers for you to initialize the relevant device pointers,
68     // which are passed to the other two functions.
69
70     // Useful snippet for error checking
71     // cudaError_t error = cudaGetLastError();
72     // if(error != cudaSuccess)
73     // {
74     //     std::cout<<"CUDA error: "<<cudaGetErrorString(error)<<std::endl;
75     //     exit(-1);
76     // }
77
78     const int Height_out = Height - K + 1;
79     const int Width_out = Width - K + 1;
80     cudaMalloc((void **) device_output_ptr, Batch*Map_out*Height_out*Width_out*sizeof(float));
81     cudaMalloc((void **) device_input_ptr, Batch*Channel*Height*Width*sizeof(float));
82     cudaMalloc((void **) device_mask_ptr, Map_out*Channel*K*K*sizeof(float));
83     cudaMemcpy(device_input_ptr, host_input, Batch*Channel*Height*Width*sizeof(float), cudaMemcpyHostToDevice);
84     cudaMemcpyToSymbol(matrix, host_mask, Map_out*Channel*K*K*sizeof(float));
85 }
86

```

## 2. Tiled shared memory convolution

- Which optimization did you choose to implement? Explain why did you choose that optimization technique.  
The next optimization I chose to implement was a tiled shared memory convolution because I thought it would allow for more hardware resources like memory to be used efficiently and enable parallelism across multiple threads at the same time.
- How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?  
This optimization works by dividing data into rectangular tiles and storing them in shared memory where convolution is carried out on all tiles simultaneously. I thought this optimization would increase performance of the forward convolution since it would reduce the bottleneck of memory bandwidth requirements by improving cache utilization and take advantage of parallelism. This optimization can synergize well with the implementation of a weighted matrix in constant memory as both techniques have a common goal of increasing memory efficiency and reducing data fetching.
- Batch profiling results:

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	.312647 ms	.675611 ms	0m 1.263s	.86
1000	2.97889 ms	6.45686 ms	0m 9.802s	.886
10000	29.624	64.7542	1m	.8714

	<i>ms</i>	<i>ms</i>	<i>36.691s</i>	
--	-----------	-----------	----------------	--

- d. Was implementing this optimization successful in improving performance? Why or why not?

*nsys Profile data:*

```
Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)
```

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
77.6	184982943	8	131238679.1	17860	564702556	cudaMemcpy
13.9	188941322	8	23505165.2	68288	193451017	cudaMalloc
7.0	95136680	6	15856100.0	3870	65019306	cudaDeviceSynchronize
1.2	16486189	6	2734364.8	16592	16289085	cudaLaunchKernel
0.2	2679811	8	334876.4	63913	865837	cudaFree

```
Generating CUDA Kernel Statistics...
Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)
```

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	94997359	2	47498679.5	29988691	65010668	conv_forward_kernel
0.0	2728	2	1368.0	1344	1376	do_not_remove_this_kernel
0.0	2634	2	1312.0	1280	1344	prefn_marker_kernel

```
CUDA Memory Operation Statistics (nanoseconds)
```

Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
91.4	955420186	2	477710053.0	391523807	563896299	[CUDA memcpy DtoH]
8.6	89529435	6	14921572.5	1216	47993284	[CUDA memcpy HtoD]

```
CUDA Memory Operation Statistics (KiB)
```


Total	Operations	Average	Minimum	Maximum	Name
1722500.0	2	861250.0	722500.000	1000000.0	[CUDA memcpy DtoH]
538919.0	6	89819.0	0.004	288906.0	[CUDA memcpy HtoD]

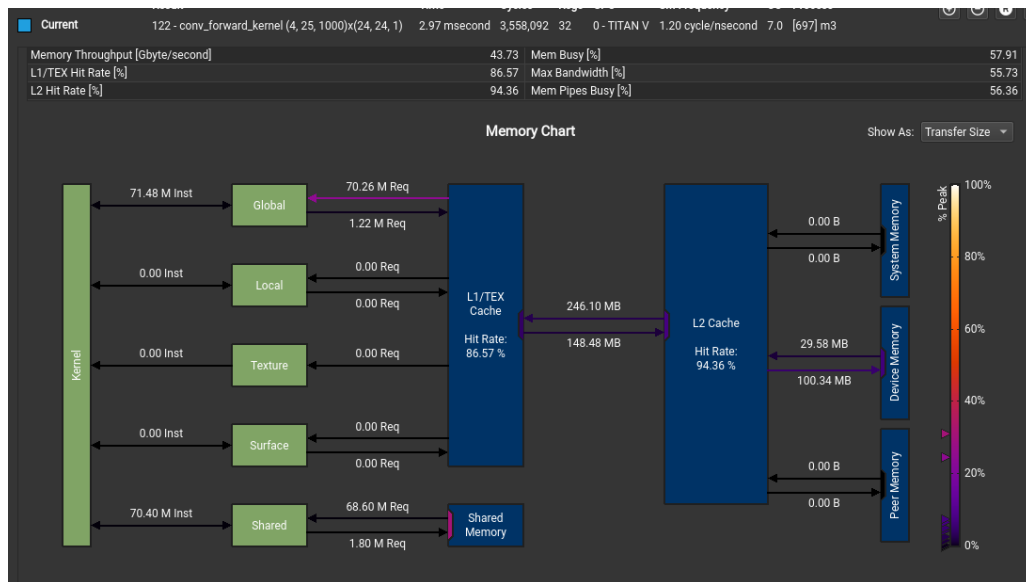
```
Generating Operating System Runtime API Statistics...
Operating System Runtime API Statistics (nanoseconds)
```

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
---------	------------	-------	---------	---------	---------	------

Based on the data displayed for the CUDA API Statistics, cudaMemcpy takes 77.6% time which is approximately 5% more usage of that API call from Milestone 2 at 71.6% and additionally the cudaDeviceSynchronize went down from 11.2% to 7.0%. This data could be attributed to utilizing more shared memory and the over 4% reduction in the cudaDeviceSynchronize could be interpreted as the code parallelizing more, thus the kernel having less of a need for that API call and the optimization being successful as also justified by the reduced OP time of approximately 95 ms.

*Nsight-Compute Profile data:*





This Nsight-Compute data shows that this optimization utilizes the same amount of SM as the weighted matrix in constant memory implementation (thus about approximately 7% more SM than Milestone 2) but actually uses less Memory at approximately 58% versus ~60% utilization for Milestone 2. This means that the tiled implementation in shared memory is a boost in performance as it makes better use of the memory bandwidth and requires less accesses.

### e. Modified code:

```

25 // Shared... float tile[...][TILE_WIDTH][TILE_WIDTH];
26 const int height_out = height - k + 1;
27 const int width_out = width - k + 1;
28 const int tile_width = TILE_WIDTH * k;
29 const int w_grid = (width_out - 1) / tile_width + 1;
30
31 // Have some nice #defs for you below to simplify indexing. Feel free to use them, or c
32 // example use of these macros:
33 // out_ad(0,0,0) = a
34 // out_ad(0,0,0) = a
35
36 #define out_ad(l1, l2, l3, l4) out_ad((l1) * (Map_out * height_out * width_out) + (l2) * (
37 #define in_ad(l1, l2, l3, l4) in_ad((l1) * (Channel * height * width) + (l2) * (height * w
38 #define mask_ad(l1, l2, l3, l4) mask_ad((l1) * (Channel * k * k) + (l2) * (k * k) + (l3) * (k
39
40 // Insert your GPU convolution kernel code here
41 int n = blockId.x;
42 int h = (blockId.y / w_grid) * tile_width + threadIdx.y;
43 int w = (blockId.y % w_grid) * tile_width + threadIdx.x;
44 int b = blockId.z;
45
46 for(int c = 0; c < Channel; c++) {
47     if(h < height_out && w < width_out) {
48         tile[c][threadIdx.y][threadIdx.x] = in_ad(b, c, h, w);
49     } else {
50         tile[c][threadIdx.y][threadIdx.x] = 0.0f;
51     }
52     __syncthreads();
53
54     float acc = 0.0f;
55     for(int c = 0; c < Channel; c++) {
56         for(int p = 0; p < k; p++) {
57             for(int q = 0; q < k; q++) {
58                 acc += tile[c][threadIdx.y][threadIdx.x] * mask_ad(n, c, p, q);
59             }
60         }
61     }
62     if(h < height_out && w < width_out) {
63         out_ad(b, n, h, w) = acc;
64     }
65 }
66
67 #undef out_ad
68 #undef in_ad
69 #undef mask_ad
70

```

## 3. Tuning with restrict and loop unrolling

- Which optimization did you choose to implement? Explain why did you choose that optimization technique.  
I chose this optimization technique because by restricting pointers and unrolling the main for loop in the kernel I thought that performance would be increased due to the reduction of memory access latency.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations? The optimization works by allowing the compiler to optimize memory accesses by assuming that memory pointed to by a restricted pointer is not aliased by any other pointer. This means this optimization would increase the performance of forward convolution since there would be improved cache utilization and reduced memory access latency. Loop unrolling would also boost performance as it would increase task parallelism. Another type of optimization that could go along with this technique is kernel fusion.

- c. Batch profiling results:

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	.231634 ms	.83974 ms	0m 1.284s	.86
1000	2.14323 ms	8.19904 ms	0m 10.354s	.886
10000	21.4202 ms	81.7411 ms	1m 36.428s	.8714

- d. Was implementing this optimization successful in improving performance? Why or why not?

*nsys Profile data:*

```
Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)
```

Time(K)	Total Time	Calls	Average	Minimum	Maximum	Name
76.8	1061431558	8	132678944.8	19519	574981321	cudaMemcpy
14.3	198136198	8	24767624.7	72197	193575437	cudaMalloc
7.5	103379586	6	17246642.7	3829	81952385	cudaDeviceSynchronize
1.2	16945637	6	2824272.8	15616	16830797	cudaLaunchKernel
0.2	2759998	8	344999.8	64750	931072	cudaFree

```
Generating CUDA Kernel Statistics...
Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)
```

Time(K)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	103452502	2	51726251.0	21581633	81958869	conv_forward_kernel
0.0	2752	2	1376.0	1344	1488	do_not_remove_this_kernel
0.0	2656	2	1328.0	1280	1376	prefn_marker_kernel

```
CUDA Memory Operation Statistics (nanoseconds)
```

Time(K)	Total Time	Operations	Average	Minimum	Maximum	Name
91.5	960951881	2	480475940.5	380771887	574180384	[CUDA memcpy DtoH]
0.5	89527364	6	14921227.3	1216	47985359	[CUDA memcpy HtoD]

```
CUDA Memory Operation Statistics (KIB)
```

Total	Operations	Average	Minimum	Maximum	Name
1722500.0	2	861250.0	722500.000	1000000.0	[CUDA memcpy DtoH]
530919.0	6	88486.5	0.004	263900.0	[CUDA memcpy HtoD]

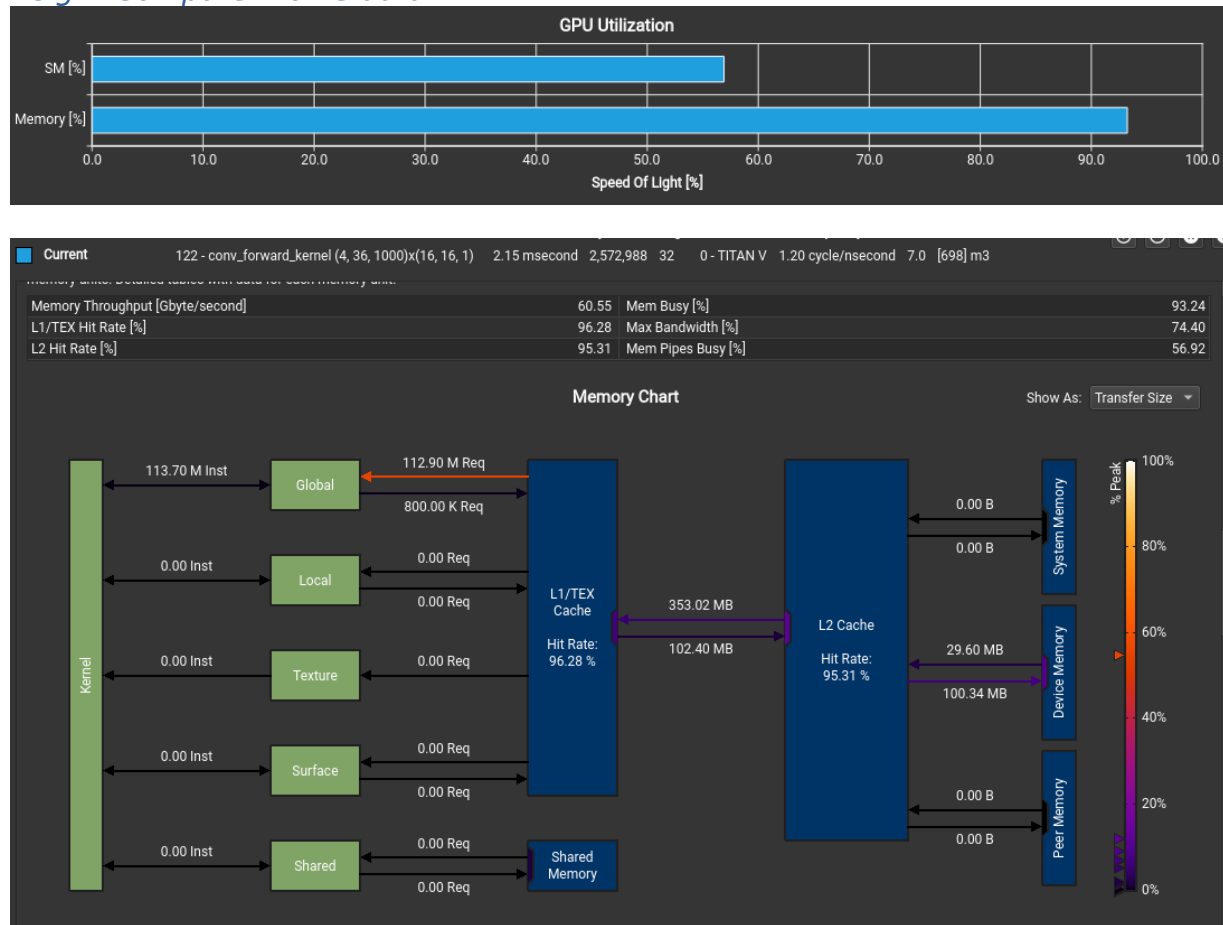
```
Generating Operating System Runtime API Statistics...
Operating System Runtime API Statistics (nanoseconds)
```

Time(K)	Total Time	Calls	Average	Minimum	Maximum	Name
---------	------------	-------	---------	---------	---------	------

According to the data above for the CUDA API calls there is over a 5% increase in cudaMemcpy from 71.6% (Milestone 2) to 76.8% and a decrease in cudaDeviceSynchronize calls from 11.2% (also Milestone 2) to 7.5%, meaning the parallelism in the code increased as there was less need for cudaDeviceSynchronize and additionally due to the reduced memory access latency and utilization of the cache the cudaMemcpy calls went up. This

optimization is thus a success in improving performance as the previously mentioned advantages helped to achieve an OP time of approximately 103 ms.

*Nsight-Compute Profile data:*



This optimization is an improvement over the baseline because we can see that uses approximately 57% of SM and 93% of memory compared to approximately 75% and 60% from Milestone 2. The L1 Cache and L2 Cache are also approximately 96% and 95% meaning that this optimization is a boost in performance as it reduces memory access latenc



e. Modified code:

```
38 // Insert your GPU convolution kernel code here
39 int m = blockIdx.x;
40 int h = (blockIdx.y / W_grid) * TILE_WIDTH + threadIdx.y;
41 int w = (blockIdx.y % W_grid) * TILE_WIDTH + threadIdx.x;
42 int b = blockIdx.z;
43
44 float acc = 0.0f;
45 for(int c = 0; c < Channel; c++) // sum over all input ch
46     acc += in_4d(b,c,h+0,w+0) * mask_4d(m,c,0,0);
47     acc += in_4d(b,c,h+0,w+1) * mask_4d(m,c,0,1);
48     acc += in_4d(b,c,h+0,w+2) * mask_4d(m,c,0,2);
49     acc += in_4d(b,c,h+0,w+3) * mask_4d(m,c,0,3);
50     acc += in_4d(b,c,h+0,w+4) * mask_4d(m,c,0,4);
51     acc += in_4d(b,c,h+0,w+5) * mask_4d(m,c,0,5);
52     acc += in_4d(b,c,h+0,w+6) * mask_4d(m,c,0,6);
53     acc += in_4d(b,c,h+1,w+0) * mask_4d(m,c,1,0);
54     acc += in_4d(b,c,h+1,w+1) * mask_4d(m,c,1,1);
55     acc += in_4d(b,c,h+1,w+2) * mask_4d(m,c,1,2);
56     acc += in_4d(b,c,h+1,w+3) * mask_4d(m,c,1,3);
57     acc += in_4d(b,c,h+1,w+4) * mask_4d(m,c,1,4);
58     acc += in_4d(b,c,h+1,w+5) * mask_4d(m,c,1,5);
59     acc += in_4d(b,c,h+1,w+6) * mask_4d(m,c,1,6);
60     acc += in_4d(b,c,h+2,w+0) * mask_4d(m,c,2,0);
61     acc += in_4d(b,c,h+2,w+1) * mask_4d(m,c,2,1);
62     acc += in_4d(b,c,h+2,w+2) * mask_4d(m,c,2,2);
63     acc += in_4d(b,c,h+2,w+3) * mask_4d(m,c,2,3);
64     acc += in_4d(b,c,h+2,w+4) * mask_4d(m,c,2,4);
65     acc += in_4d(b,c,h+2,w+5) * mask_4d(m,c,2,5);
66     acc += in_4d(b,c,h+2,w+6) * mask_4d(m,c,2,6);
67     acc += in_4d(b,c,h+3,w+0) * mask_4d(m,c,3,0);
68     acc += in_4d(b,c,h+3,w+1) * mask_4d(m,c,3,1);
69     acc += in_4d(b,c,h+3,w+2) * mask_4d(m,c,3,2);
70     acc += in_4d(b,c,h+3,w+3) * mask_4d(m,c,3,3);
71     acc += in_4d(b,c,h+3,w+4) * mask_4d(m,c,3,4);
72     acc += in_4d(b,c,h+3,w+5) * mask_4d(m,c,3,5);
73     acc += in_4d(b,c,h+3,w+6) * mask_4d(m,c,3,6);
74     acc += in_4d(b,c,h+4,w+0) * mask_4d(m,c,4,0);
75     acc += in_4d(b,c,h+4,w+1) * mask_4d(m,c,4,1);
76     acc += in_4d(b,c,h+4,w+2) * mask_4d(m,c,4,2);
77     acc += in_4d(b,c,h+4,w+3) * mask_4d(m,c,4,3);
78     acc += in_4d(b,c,h+4,w+4) * mask_4d(m,c,4,4);
79     acc += in_4d(b,c,h+4,w+5) * mask_4d(m,c,4,5);
80     acc += in_4d(b,c,h+4,w+6) * mask_4d(m,c,4,6);
81     acc += in_4d(b,c,h+5,w+0) * mask_4d(m,c,5,0);
82     acc += in_4d(b,c,h+5,w+1) * mask_4d(m,c,5,1);
```

#### 4. Using Streams to overlap computation with data transfer

- a. Which optimization did you choose to implement? Explain why did you choose that optimization technique.

I chose the optimization technique of implementing streams to overlap computation with data transfer because I thought it would make the code more seamless by reducing the bottleneck of time required for data transferred and that the streams would synergize with each other in order to further the performance.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?
- This optimization works by dividing the device resources into different streams which are all tasked with different operations such as data transfer or convolution. These streams can work hand in hand with one another to help one another complete their tasks and this significantly cuts down on the time needed for data transfer as it

is done concurrently with the convolution. As a result of the multifaceted capabilities of streams, the performance is enhanced by cutting down on time and increasing parallelism. This optimization technique can synergize with techniques like memory pooling or batching to reuse device memory for operations.

c. Batch profiling results:

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	.000685 ms	.001195 ms	0m 1.231s	.86
1000	.000758 ms	.000903 ms	0m 11.159s	.886
10000	.00071 ms	.001263 ms	1m 48.981s	.8714

d. Was implementing this optimization successful in improving performance? Why or why not?

```
Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)
Time(K) Total Time Calls Average Minimum Maximum Name
-----
86.2 119080802 42 2848088.8 11159 8764623 cudaMemcpyAsync
12.6 174369767 8 21766220.9 8211 172181476 cudaMemcpy
1.0 13194216 24 549319.0 22864 12488116 cudaLaunchKernel
0.2 3261644 2 1631921.0 19389 3144255 cudaMemcpy
0.0 141134 0 17641.8 482 93238 cudaFree
0.0 124644 20 6232.2 1347 72818 cudaStreamCreate
0.0 122185 20 6189.3 1844 53821 cudaStreamDestroy
0.0 32192 8 4161.3 1683 18225 cudaDeviceSynchronize

Generating CUDA Kernel Statistics...
Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)
Time(K) Total Time Instances Average Minimum Maximum Name
-----
100.0 78493171 20 3924658.5 1617299 6271487 conv_forward_kernel
0.0 2752 2 1376.0 1376 1376 prefn_marker_kernel
0.0 2496 2 1248.0 1248 1248 do_not_remove_this_kernel


CUDA Memory Operation Statistics (nanoseconds)
Time(K) Total Time Operations Average Minimum Maximum Name
-----
93.1 1818006296 20 90900314.8 39822220 64926271 [CUDA memcpy DtoH]
0.9 75619790 24 3125824.6 1151 4588956 [CUDA memcpy HtoD]

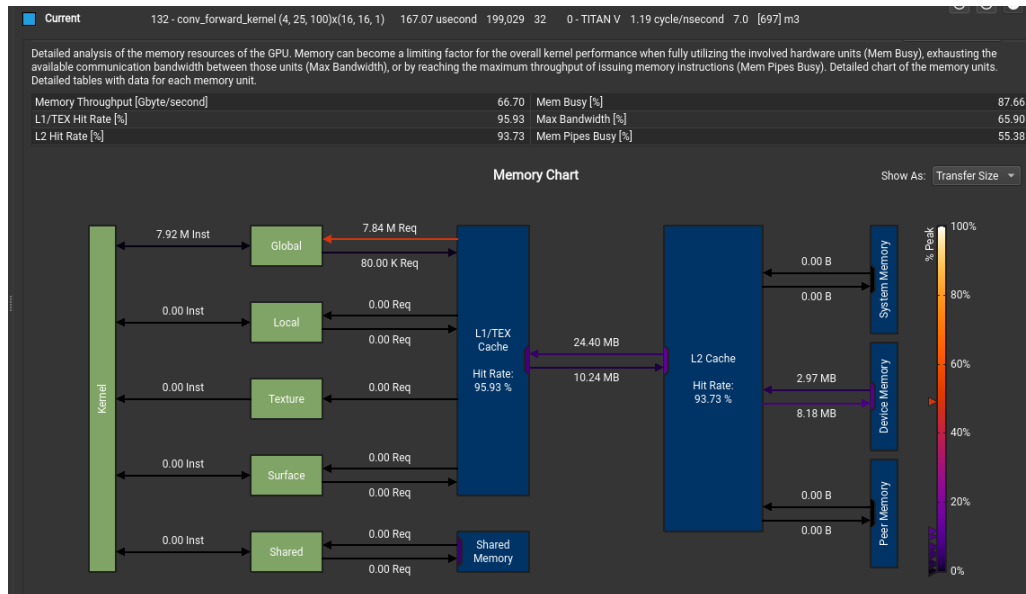
CUDA Memory Operation Statistics (KiB)
Total Operations Average Minimum Maximum Name
-----
1722588.0 20 86125.0 72258.860 188888.0 [CUDA memcpy DtoH]
538919.0 24 22454.0 0.004 28898.0 [CUDA memcpy HtoD]

Generating Operating System Runtime API Statistics...
```

Based on the CUDA API Statistics it seems that clearly cudaMemcpyAsync is called the most for 86.2% of the time compared to cudaMemcpy being called for 71.6% of the time in Milestone 2. This change can be attributed to the fact that we are utilizing streams to simultaneously transfer data and compute the convolution, vastly enhancing the performance time. The only potential cause for concern though is that the OP times are in the ~.001 ms for this optimization method meaning that although this was successful we are unsure as to what extent the performance was enhanced.

Nsight-Compute Profile data:





For this implementation of stream optimization, the SM and Memory usages are very high because the streams handle both the convolution and data transfer of the kernel, meaning that the code is an improvement due to it utilizing the hardware resources of the GPU to their fullest to achieve the kernel's task. The only underlying issue is still the fact that the OP times are somewhat suspect meaning that there must be underlying data points that need to be measured that I have not found on Nsight\_Compute.

#### e. Modified code:

*In conv\_forward\_gpu\_prolog:*

```

80 const int Height_out = Height - K + 1;
81 const int Width_out = Width - K + 1;
82 float* temp_host_output = (float*)host_output;
83 int STREAM_X = (Batch*Channel*Height*Width)/10;
84 int STREAM_Y = (Batch*Map_out*Height_out*Width_out)/10;
85 int W_grd = (Width_out+TILE_WIDTH-1)/TILE_WIDTH;
86 int H_grd = (Height_out+TILE_WIDTH-1)/TILE_WIDTH;
87 int Y_grd = W_grd * H_grd;
88 dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1);
89 dim3 gridDim(Map_out, Y_grd, Batch/10);
90 cudaMalloc((void**) &device_output_ptr, Batch*Map_out*Height_out*Width_out*sizeof(float));
91 cudaMalloc((void**) &device_input_ptr, Batch*Channel*Height*Width*sizeof(float));
92 cudaMalloc((void**) &device_mask_ptr, Map_out*Channel*K*K*sizeof(float));
93
94 cudaStream_t stream[10];
95 for(int i = 0; i < 10; i++)
96     cudaStreamCreate(&stream[i]);
97 cudaMemcpyAsync(&device_mask_ptr, host_mask, Map_out*Channel*K*K*sizeof(float), cudaMemcpyHostToDevice, stream[0]);
98 for(int i = 0; i < 10; i++)
99     {
100         int x = STREAM_X*i;
101         int y = STREAM_Y*i;
102         cudaMemcpyAsync(&device_input_ptr+x, host_input+x, STREAM_X*sizeof(float), cudaMemcpyHostToDevice, stream[i]);
103         conv_forward_kernel<<gridDim, blockDim, 0, stream[i]>>(&device_output_ptr+y, &device_input_ptr+x, &device_mask_ptr, Batch, Map_out, Channel, Height, Width, K);
104         cudaMemcpyAsync(temp_host_output+y, &device_output_ptr+y, STREAM_Y*sizeof(float), cudaMemcpyDeviceToHost, stream[i]);
105     }
106 cudaDeviceSynchronize();
107 cudaStreamDestroy(stream[0]);
108 cudaStreamDestroy(stream[1]);
109 cudaStreamDestroy(stream[2]);
110 cudaStreamDestroy(stream[3]);
111 cudaStreamDestroy(stream[4]);
112 cudaStreamDestroy(stream[5]);
113 cudaStreamDestroy(stream[6]);
114 cudaStreamDestroy(stream[7]);
115 cudaStreamDestroy(stream[8]);
116 cudaStreamDestroy(stream[9]);
117 cudaFree(device_output_ptr);
118 cudaFree(device_input_ptr);
119 cudaFree(device_mask_ptr);
120

```